

元学习——从MAML到MAML++

作者：凯鲁嘎吉 - 博客园 <http://www.cnblogs.com/kailugaji/>

Few-shot learning领域最近有了实质性的进展。这些进步大多来自于将few-shot learning作为元学习问题。[Model-Agnostic Meta-Learning](#) (MAML)是目前利用元学习进行few-shot learning的最佳方法之一。MAML简单，优雅，功能强大，但是它有很多问题，比如对神经网络结构非常敏感，经常导致训练时不稳定，需要费力的超参数搜索来稳定训练和实现高泛化，并且在训练和推理时间上都非常昂贵的计算。在文"[How to train your MAML](#)"中，对MAML进行了各种改进，不仅稳定了系统，而且大幅度提高了MAML的泛化性能、收敛速度和计算开销。所提方法称之为MAML++。本博文首先介绍什么是元学习，经典的Model-Agnostic Meta-Learning的定义与执行过程，进而说明MAML面临的缺点与挑战，针对这些问题，进行相应改进，从而得到MAML++。

1. Meta Learning (Learn to Learn)

How to train your MAML

➤ Meta Learning

Supervised Learning:

Inputs: \mathbf{x} Outputs: \mathbf{y}
 $y = f(\mathbf{x}; \theta)$

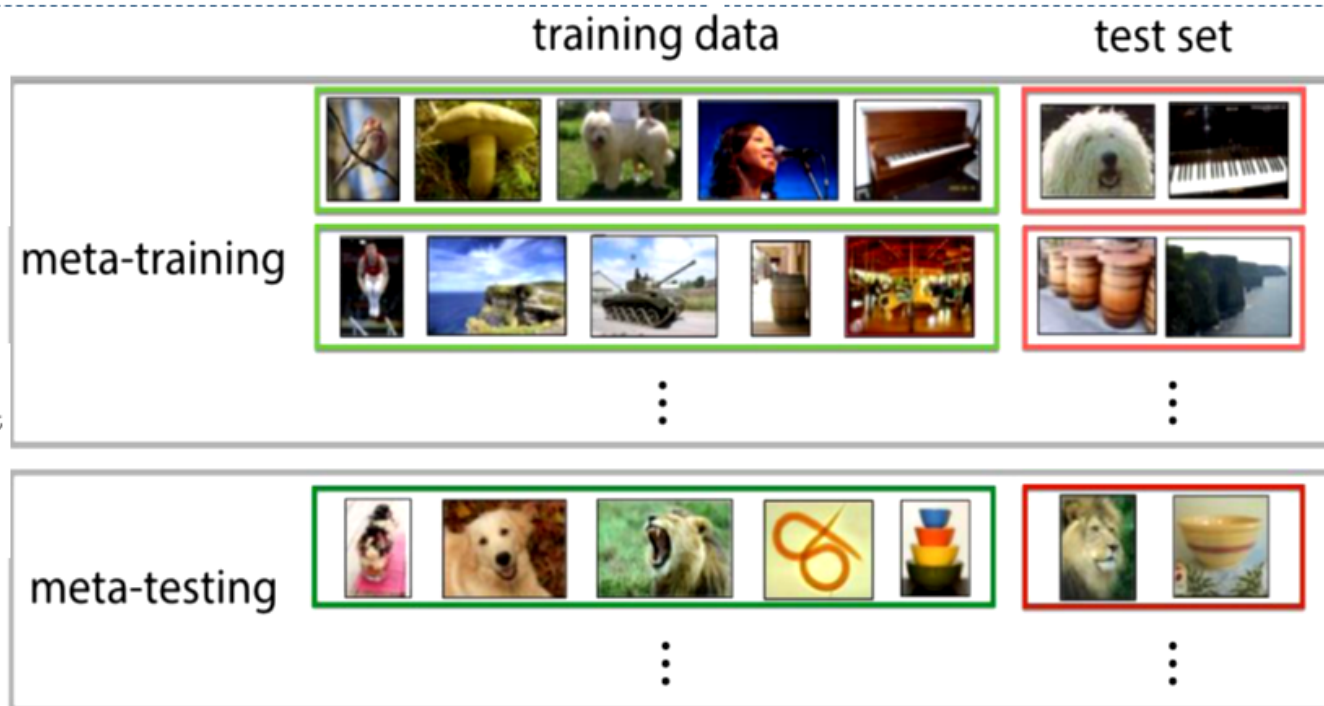
Data: $\{(\mathbf{x}, \mathbf{y})_i\}$

Meta-Supervised Learning:

Inputs: $\mathcal{D}_{\text{train}}$ \mathbf{x}_{test} Outputs: \mathbf{y}_{test}
 $\{(\mathbf{x}, \mathbf{y})_{1:K}\}$
 $\mathbf{y}_{\text{test}} = f(\mathcal{D}_{\text{train}}, \mathbf{x}_{\text{test}}; \theta)$

Data: $\{\mathcal{D}_i\}$

$\mathcal{D}_i : \{(\mathbf{x}, \mathbf{y})_j\}$




➤ Why Learn to Learn?

- effectively **reuse data** on other tasks
- **replace manual engineering** of architecture, hyperparameters, etc.
- learn to **quickly adapt to unexpected scenarios** (inevitable failures, long tail)
- learn how to learn **with weak supervision**

Finn, C., Abbeel, P. & Levine, S. Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks. ICML 2017.

2. Black-Box Adaption vs Optimization-Based Approach

➤ Black-Box Adaptation




1. Sample task \mathcal{T}_i (or mini batch of tasks)
2. Sample disjoint datasets $\mathcal{D}_i^{\text{tr}}, \mathcal{D}_i^{\text{test}}$ from \mathcal{D}_i
3. Compute $\phi_i \leftarrow f_\theta(\mathcal{D}_i^{\text{tr}})$
4. Update θ using $\nabla_\theta \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{test}})$

➤ Optimization-Based Approach

$$\max_{\phi_i} \log p(\mathcal{D}_i^{\text{tr}} | \phi_i) + \log p(\phi_i | \theta)$$

$$y^{\text{ts}} = f_{\text{MAML}}(\mathcal{D}_i^{\text{tr}}, x^{\text{ts}}) = f_{\phi_i}(x^{\text{ts}})$$

$$\text{where } \phi_i = \theta - \alpha \nabla_\theta \mathcal{L}(\theta, \mathcal{D}_i^{\text{tr}})$$



1. Sample task \mathcal{T}_i (or mini batch of tasks)
2. Sample disjoint datasets $\mathcal{D}_i^{\text{tr}}, \mathcal{D}_i^{\text{test}}$ from \mathcal{D}_i
3. ~~Compute $\phi_i \leftarrow f_\theta(\mathcal{D}_i^{\text{tr}})$~~ Optimize $\phi_i \leftarrow \theta - \alpha \nabla_\theta \mathcal{L}(\theta, \mathcal{D}_i^{\text{tr}})$
4. Update θ using $\nabla_\theta \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{test}})$

➤ MAML

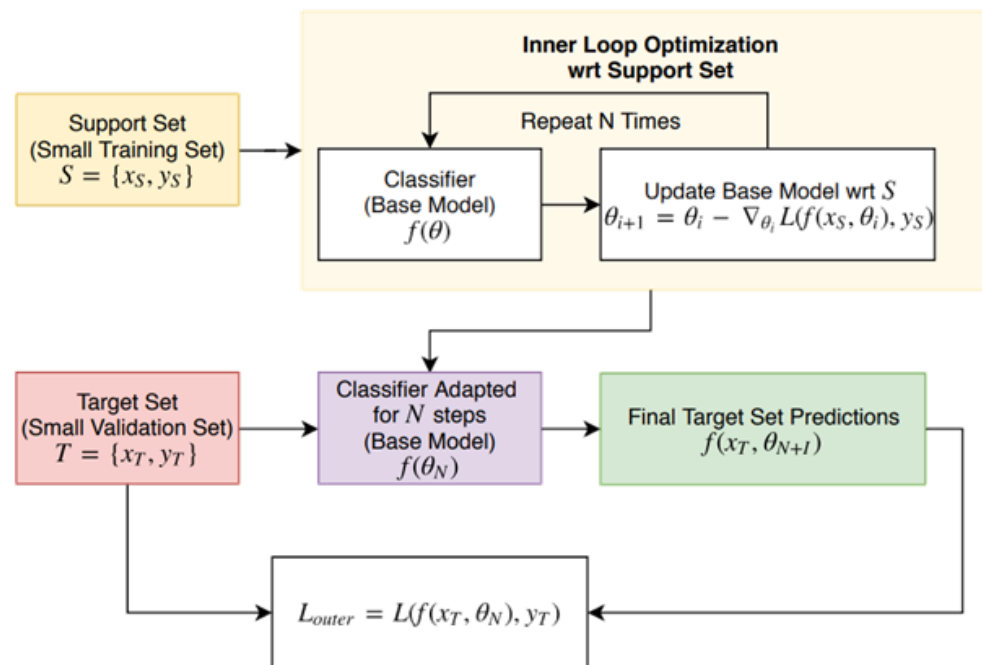
Meta-Train

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

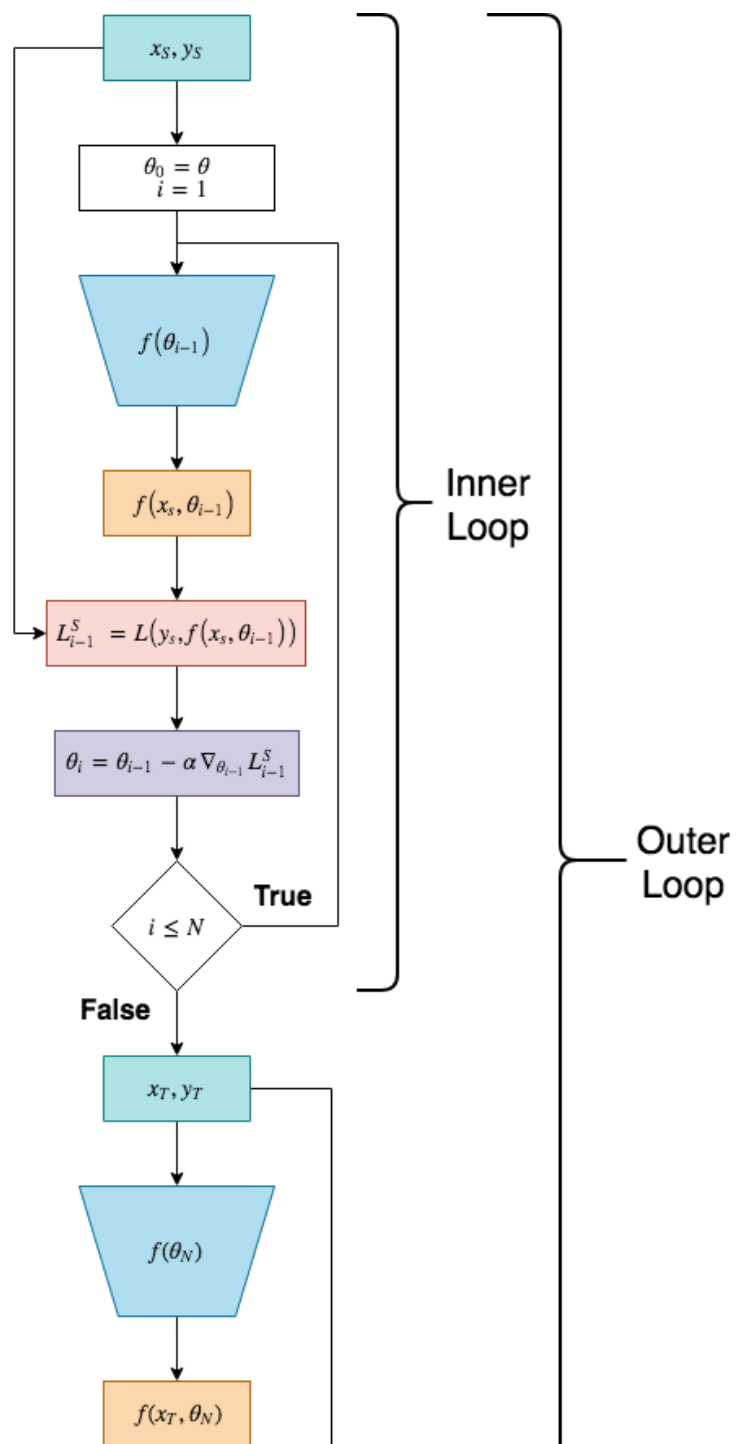
```
1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples
6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  Support Set
7:   end for
8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$  Query Set
9: end while
```



- ❑ 无论是meta-train还是fine-tune阶段，每个task都包括两部分，support set和query set。
- ❑ 不同的是，在meta-train阶段，support set参与第一次参数更新。这里的参数更新并没有直接作用于原模型，我们可以理解为先copy了一下模型，用来计算新参数。利用第一轮更新后的参数，通过query set计算第二轮梯度，这一轮的梯度才是模型真正用于更新参数的梯度；
- ❑ 在fine-tune阶段，support set参与第一次参数更新，更新结果直接作用于原模型，此时没有第二次参数更新，因为query set相当于测试集。
- ❑ 作者：徐不知，链接：<https://www.zhihu.com/question/292959709/answer/605504088>

Finn, C., Abbeel, P. & Levine, S. Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks. ICML 2017.

3




4. MAML Problems

$$L_N^T = L(y_T, f(x_s, \theta_N))$$

How to train your MAML

➤ MAML->MAML++

$$\mathbf{y}_{\text{test}} = f(\mathbf{x}_{\text{test}}; \theta - \alpha \nabla_{\theta} \mathcal{L}(\mathcal{D}_{\text{train}}))$$



MAML Problems:

- Training Instability
- Second Order Derivative Cost
- Absence of Batch Normalization Statistic Accumulation
- Shared (across step) Batch Normalization Bias
- Shared Inner Loop (across step and across parameter) Learning Rate
- Fixed Outer Loop Learning Rate

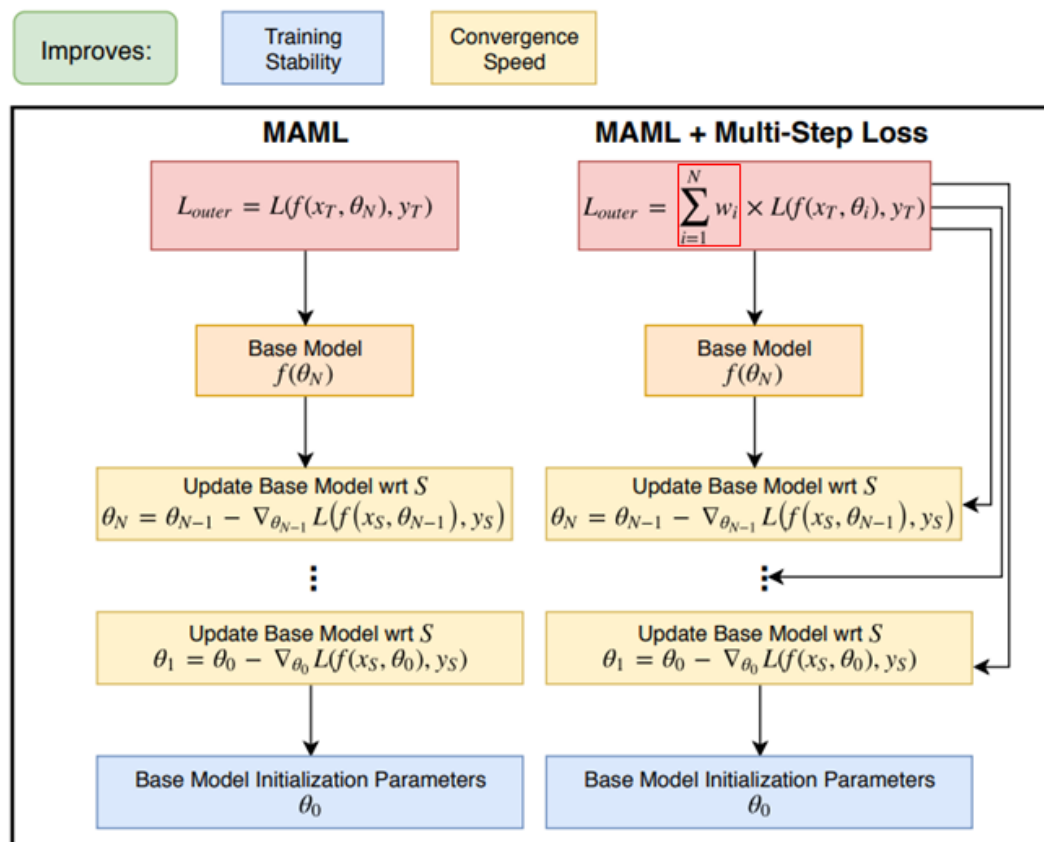
Antoniou, A., Edwards, H., & Storkey, A. How to train your MAML. ICLR 2019.

4

5. MAML++

➤ MAML->MAML++

- Gradient Instability->**Multi-Step** Loss Optimization (MSL)



- Minimizing the target set loss computed by the base-network after every step towards a support set task.
- More specifically, the loss minimized is a weighted sum of the target set losses after every support set loss update.

$$\theta = \theta - \beta \nabla_{\theta} \sum_{b=1}^B \sum_{i=0}^N v_i \mathcal{L}_{T_b}(f_{\theta_i^b})$$

task b , step i

➤ MAML->MAML++

- Absence of Batch Normalization Statistic Accumulation->**Per-Step** Batch Normalization Running Statistics (BNRS)
- Shared (across step) Batch Normalization Bias->**Per-Step** Batch Normalization Weights and Biases (BNWB)

Improves:

Training
Stability

Convergence
Speed

Generalization
Performance

Given Input x at step i with shape (b, f, h, w) and N inner loop optimization steps

MAML Batch Normalization

Running Mean

μ
shape: (f)

Scaling Parameters

γ
shape: (f)

Running Std Deviation

σ
shape: (f)

Shift Parameters

β
shape: (f)

$$bn(x, i) = \beta + \gamma \left(\frac{x - \mu}{\sigma} \right)$$

Per-Step Batch Normalization

Running Mean

μ
shape: (N, f)

Scaling Parameters

γ
shape: (N, f)

Running Std Deviation

σ
shape: (N, f)

Shift Parameters

β
shape: (N, f)

$$bn(x, i) = \beta_i + \gamma_i \left(\frac{x - \mu_i}{\sigma_i} \right)$$

- Collecting statistics in a per-step regime. Instantiating N sets of running mean and running standard deviation for each batch normalization layer in the network and update the running statistics respectively with the steps being taken during the optimization.
- Learning a set of biases per-step within the inner-loop update process.

➤ MAML->MAML++

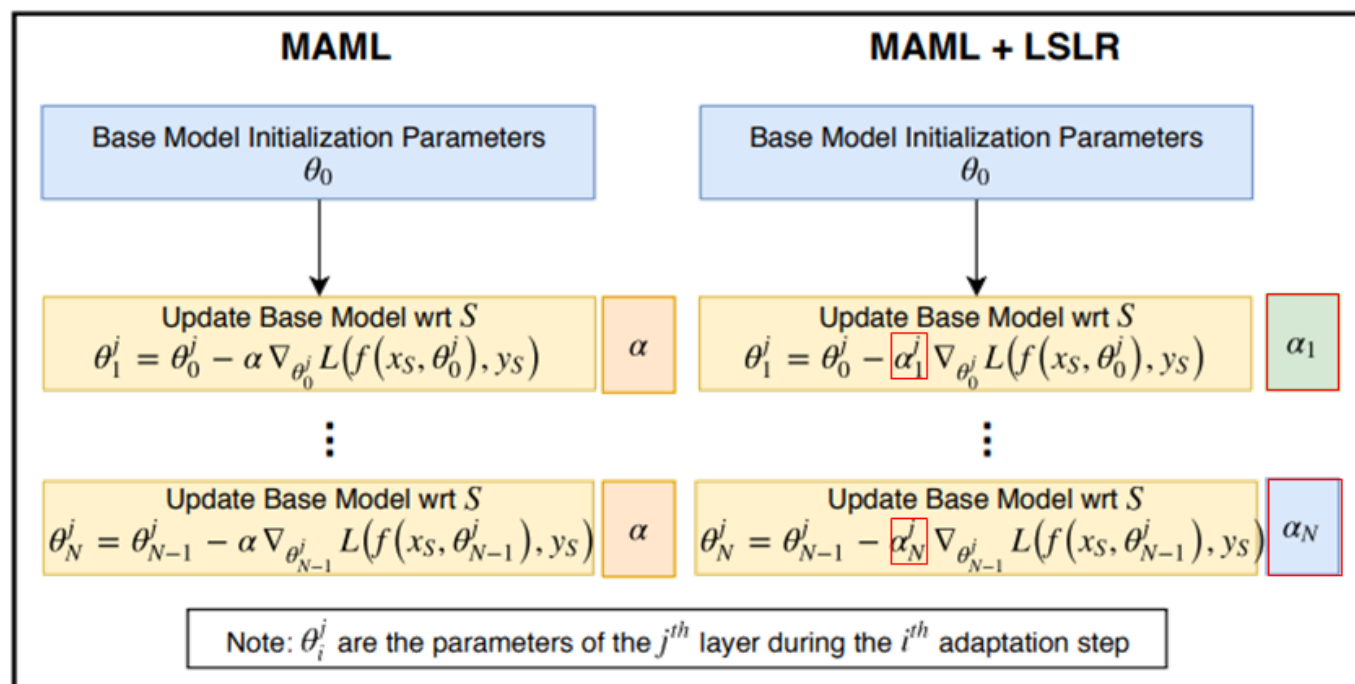
- Shared Inner Loop (across step and across parameter) Learning Rate->Learning **Per-Layer Per-Step** Learning Rates and Gradient Directions (LSLR)

Improves:

Training
Stability

Convergence
Speed

Generalization
Performance



- Learning a learning rate and direction for each layer in the network.
- Learning different learning rates for each adaptation of the base-network.

How to train your MAML

➤ MAML->MAML++


- Second Order Derivative Cost->Derivative-Order Annealing (DA)
 - ❑ Using first-order gradients for the first 50 epochs of the training phase, then switching to second-order gradients for the remainder of the training phase.
- Fixed Outer Loop Learning Rate->Cosine Annealing of Meta-Optimizer Learning Rate (CA)
 - ❑ Applying the cosine annealing scheduling on the meta-model's optimizer (i.e. the meta-optimizer).

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

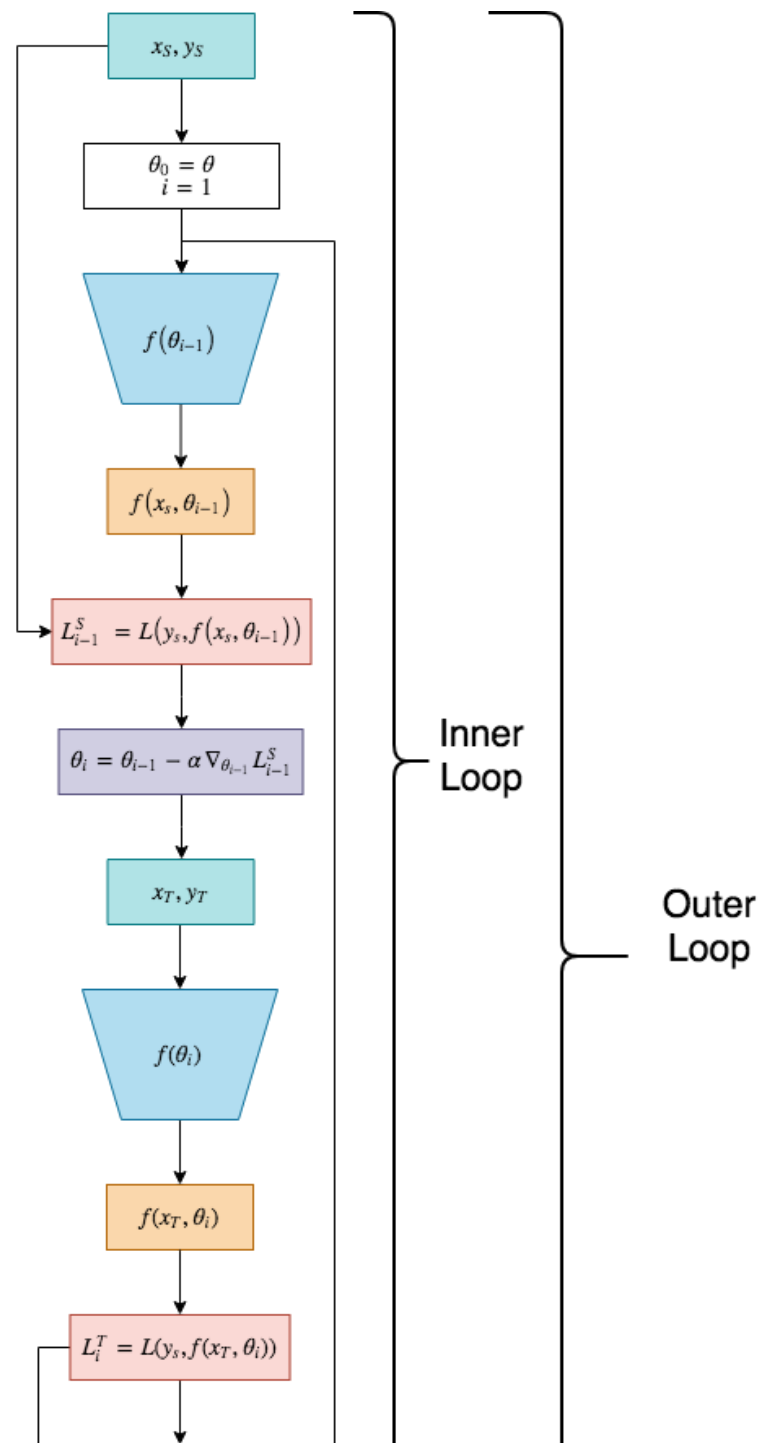
Approach	Accuracy			
	Omniglot 20-way		Mini-ImageNet 5-way	
	1-shot	5-shot	1-shot	5-shot
Siamese Nets	88.2%	97.0%	-	-
Matching Nets	93.8%	98.5%	43.56%	55.31%
Neural Statistician	93.2%	98.1%	-	-
Memory Mod.	95.0%	98.6%	-	-
Meta-SGD	95.93±0.38%	98.97±0.19%	50.47±1.87%	64.03±0.94%
Meta-Networks	97.00%	—	49.21%	-
MAML (original)	95.8±0.3%	98.9±0.2%	48.70±1.84%	63.11±0.92%
MAML (local replication)	91.27±1.07%	98.78%	48.25±0.62%	64.39±0.31%
MAML++	97.65±0.05%	99.33±0.03%	52.15±0.26%	68.32±0.44%

MAML++ Few-Shot Results

Antoniou, A., Edwards, H., & Storkey, A. How to train your MAML. ICLR 2019.



- ❑ The results of the approach indicate that learning per-step learning rates, batch normalization parameters and optimizing on per-step target losses appears to be key for fast, highly automatic and strongly generalizable few-shot learning.



6. 参考文献

[1] Finn, C., Abbeel, P. & Levine, S. [Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks](#). ICML 2017.

Code: <https://github.com/cbfinn/maml>, <https://github.com/dragen1860/MAML-Pytorch>

Finn个人主页: <https://ai.stanford.edu/~cbfinn/>

[2] Antoniou, A., Edwards, H., & Storkey, A. [How to train your MAML](#). ICLR 2019. Code: <https://github.com/AntreasAntoniou/HowToTrainYourMAMLPytorch>

[3] How to train your MAML: A step by step approach · BayesWatch <https://www.bayeswatch.com/2018/11/30/HTYM/>

[4] CS 330: Deep Multi-Task and Meta Learning <http://web.stanford.edu/class/cs330/>

[5] Meta-Learning: Learning to Learn Fast <https://lilianweng.github.io/lil-log/2018/11/30/meta-learning.html>

